

# Problem1

Pejman

Nov 28th 2018

## 1 Introduction

### 1.1 Question:

Use freely available data from the web to predict/explain macroeconomic indicators. Financial/Economic/Fundamentals data are not allowed.

### 1.2 Description of the Approach and Steps

*Description.-* In the following we have gathered data from the internet, from an open source. These data includes GDP historical data from 1941 and a few other time series (historical data) of (potentially) indicators such as the US population, durables, and 3-month LIBOR.

*Objective.-* The objective is to find which of indicators above, or combination of them is a good proxy to the GDP.

As a second goal and second phase we may attempt to forecast the GDP using these indicators.

*Steps.-* Steps taken are as follows:

- Using <https://fred.stlouisfed.org/> as the source, we gathered data of the following indicators: Population, Libor, Durables, GDP
- Since the GDP is quarterly reported an appropriate averaging over every three months of the other data is made.
- The start and end dates are different and filtering was required.
- A roll of one day is performed to match the dates with the GDP report dates
- Illustration of the normalized data. Normalization is done by subtracting the mean and division by the max-min in the time series window.
- (although visually observable) The  $r^2$ -squared is performed by fitting a line in the scatter plot of the indicator-GDP.
- We found the population to be in the best agreement with the GDP.
- Extension...(if time allows) forward propagation of the indicators and prediction of the GDP.

## 2 Packages

```
In [9]: '''
        Source of data : https://fred.stlouisfed.org/
        Author : Pej
        '''

import pandas as pd
import datetime
import matplotlib.pyplot as plt

### other packs
from sklearn import linear_model
```

## 3 Cleaning and Acquiring Data

```
In [15]: ### reading data and cleaning (less needed here though)
        #####
        ##### Actual GDP
        #####
        gdp_data = pd.read_csv('./GDPC1.csv', index_col=0, parse_dates=[0])
        gdp_data.columns = ['Actual GDP']
        gdp_data = gdp_data.reset_index()
        gdp_data['DATE'] = pd.to_datetime(gdp_data['DATE'])
        gdp_data = gdp_data.set_index('DATE')

        #####
        ##### Oil Production (Finished_Motor_Gasoline)
        #####
        ##### using the data on oil production from 1991 to 2018
        oil_production = \
            pd.read_csv('./Weekly_US_Product_Supplied_of_Finished_Motor_Gasoline.csv',
                        skiprows=[0,1,2, 3])
        oil_production.columns=['Date', 'Productions']
        oil_production.loc[:, 'Date'] = pd.to_datetime(oil_production['Date'])
        oil_production_cleaned = oil_production.set_index('Date')
        ### resampling Quarterly
        oil_prod_quarterly = oil_production_cleaned.resample('Q').mean()

        #####
        ##### 3-MONTH LIBOR (some cleaning...)
        #####
        libor3M = pd.read_csv('./USD3MTD156N.csv')
        libor3M['DATE'] = pd.to_datetime(libor3M['DATE'])
        libor3M = libor3M.set_index('DATE')
```

```

libor3M['USD3MTD156N'] = pd.to_numeric(libor3M['USD3MTD156N'].values[:],
                                         errors='coerce')
null_col = libor3M.columns[libor3M.isnull().any()] # obvious !
num_nulls = len(libor3M[libor3M.isnull().any(axis=1)][null_col]) ## ne 0
libor3M = libor3M.dropna() # num_nulls = 0

### This is a three month libor but in case the daycount is different...
libor3M = libor3M.resample('Q').mean()

#####
#### DURABLE GOODS ORDERS
#####
durable = pd.read_csv('./DGORDER.csv')
durable = durable.dropna(how='any')
durable['DATE'] = pd.to_datetime(durable['DATE'])
durable = durable.set_index('DATE')
### resampling
durable = durable.resample('Q').mean()

#####
#### US Population
#####
population = pd.read_csv('./POPTHM.csv')
population = population.dropna(how='any')
population['DATE'] = pd.to_datetime(population['DATE'])
population = population.set_index('DATE')
### resampling
population = population.resample('Q').mean()

```

## 4 Illustration

```

In [16]: #####
### finding the row number of the most recent date
### there might be better ways ...
#####

t0_oil   = datetime.datetime.strptime(str(oil_prod_quarterly.index[0]), "%Y-%m-%d %H:%M:%S")
t0_gdp   = datetime.datetime.strptime(str(gdp_data.index[0]), "%Y-%m-%d %H:%M:%S")
t0_libor = datetime.datetime.strptime(str(libor3M.index[0]), "%Y-%m-%d %H:%M:%S")
t0_pop   = datetime.datetime.strptime(str(population.index[0]), "%Y-%m-%d %H:%M:%S")
t0_durb  = datetime.datetime.strptime(str(durable.index[0]), "%Y-%m-%d %H:%M:%S")

tf_oil   = datetime.datetime.strptime(str(oil_prod_quarterly.index[-1]), "%Y-%m-%d %H:%M:%S")
tf_gdp   = datetime.datetime.strptime(str(gdp_data.index[-1]), "%Y-%m-%d %H:%M:%S")

```

```

tf_libor = datetime.datetime.strptime(str(libor3M.index[-1]), "%Y-%m-%d %H:%M:%S")
tf_pop    = datetime.datetime.strptime(str(population.index[-1]), "%Y-%m-%d %H:%M:%S")
tf_durb   = datetime.datetime.strptime(str(durable.index[-1]), "%Y-%m-%d %H:%M:%S")

start_date = max([t0_oil, t0_gdp, t0_libor, t0_pop, t0_durb ])
end_date    = min([tf_oil, tf_gdp, tf_libor, tf_pop, tf_durb ])

##### filter and adjust the date
libor3M = libor3M.loc[start_date:end_date]
libor3M = libor3M.reset_index()
libor3M['DATE'] =\
    libor3M['DATE'].apply(lambda x: x + datetime.timedelta(days=1))
libor3M = libor3M.set_index('DATE')

oil_prod_quarterly = oil_prod_quarterly.loc[start_date:end_date]
oil_prod_quarterly = oil_prod_quarterly.reset_index()
oil_prod_quarterly['Date'] =\
    oil_prod_quarterly['Date'].apply(lambda x: x + datetime.timedelta(days=1))
oil_prod_quarterly = oil_prod_quarterly.set_index('Date')

population = population.loc[start_date:end_date]
population = population.reset_index()
population['DATE'] =\
    population['DATE'].apply(lambda x: x + datetime.timedelta(days=1))
population = population.set_index('DATE')

durable = durable.loc[start_date:end_date]
durable = durable.reset_index()
durable['DATE'] =\
    durable['DATE'].apply(lambda x: x + datetime.timedelta(days=1))
durable = durable.set_index('DATE')

```

## 5 Analysis

### 5.1 R2-Squared

```

In [17]: ##### visualization:
final_dt = pd.concat([oil_prod_quarterly,
                      libor3M,
                      population,
                      durable,
                      gdp_data], axis=1).dropna()

final_dt.columns=['oil',
                  'libor',

```

```

        'population',
        'durable',
        'GDP',
    ]

    ## normalize for illustration
    final_dt_norm = (final_dt - final_dt.mean())/(final_dt.max() - final_dt.min())


    ### seems like labor is not a good fit unless we consider some non-linear relation
    ### between Labor and GDP ... so let's drop it for now.
    fig, ax = plt.subplots(figsize=[10,10])
    for col in range(final_dt_norm.shape[1]):
        if final_dt_norm.columns[col] != 'labor':
            ax.plot(final_dt_norm.values[:, col],
                    label=str(final_dt_norm.columns[col])+'_normalized')
    ax.set_xlabel('quarters since 1992', **{'fontsize':18})
    plt.legend()
    plt.show()

    ##### let's now the correlation variace around the fit
    ##### with the gdp or not.
    ##### From the plot they are not perfectly in shape ... but anyway...

    y_data = final_dt.iloc[:, -1]          ## gdp the target

    for i in range(4):    ### since we have chosen four indicators
        X_data = final_dt.iloc[:, [i]] ## the indexes we chose
        ##
        lin_mod = linear_model.LinearRegression()
        _ = lin_mod.fit(X_data, y_data)
        print ' ===== '
        print 'The indecator chosen is : %s ' %(final_dt.columns[i])
        print('R-Squared of Linear Regression Model:', lin_mod.score(X_data, y_data))
        print ' ===== \n\n'

```



```
=====
The indecator chosen is : oil
('R-Squared of Linear Regression Model:', 0.7433326914980255)
=====
```

```
=====
The indecator chosen is : labor
('R-Squared of Linear Regression Model:', 0.44945552780276155)
=====
```

```
=====
```

```
The indecator chosen is : population
('R-Squared of Linear Regression Model:', 0.98341604629017587)
```

```
=====
```

```
=====
```

```
The indecator chosen is : durable
('R-Squared of Linear Regression Model:', 0.78012867601848335)
```

```
=====
```

## 5.2 Others